

Computer Architecture & Real-Time Operating System

7. Processor Architecture (2/2) (Microarchitecture)

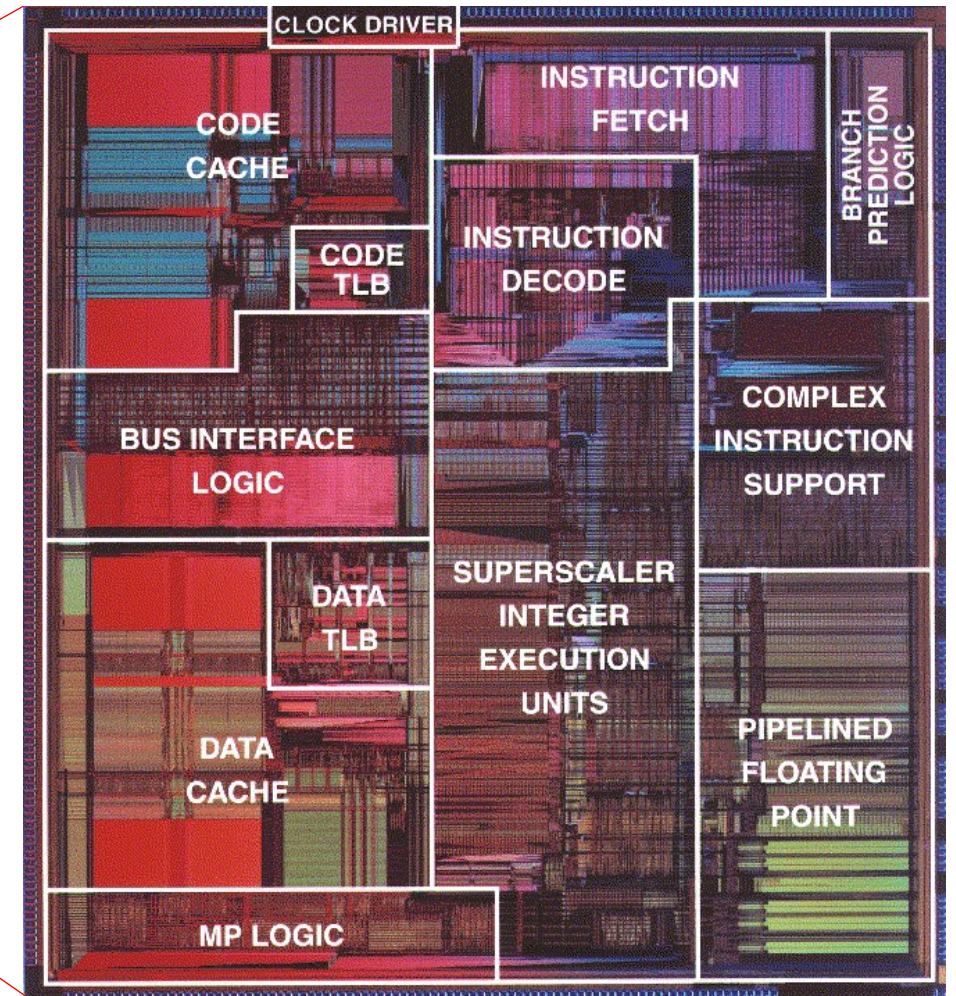

Prof. Jong-Chan Kim

Dept. Automobile and IT Convergence



Microarchitecture


- What's happening inside a chip?




Die photo of Intel Pentium Processor

Moore's Law

- Number of **transistors** in a CPU doubles every 24 months




Economist.com




Gordon Moore (Co-founder of Intel)

Transistor




Source: <https://ko.wikipedia.org/wiki/트랜지스터>


From Transistors to CPU



Transistor




Circuit



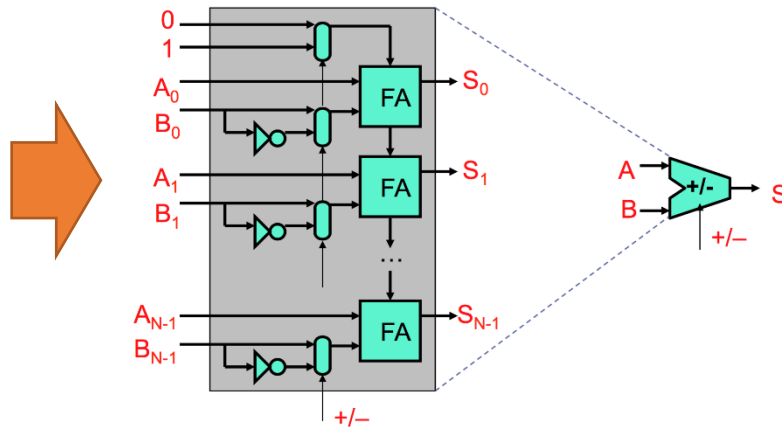
Logic Gate

Full Adder

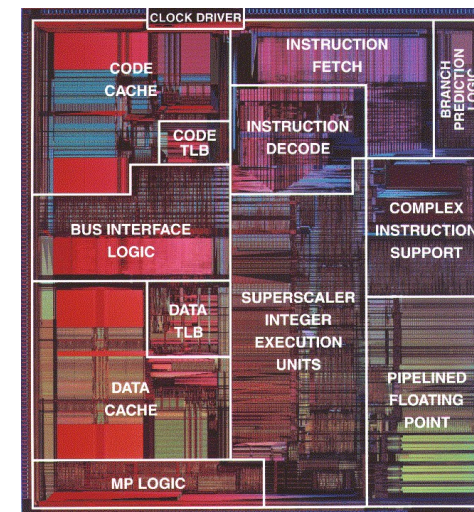


Module

N-bit Adder/Subtractor

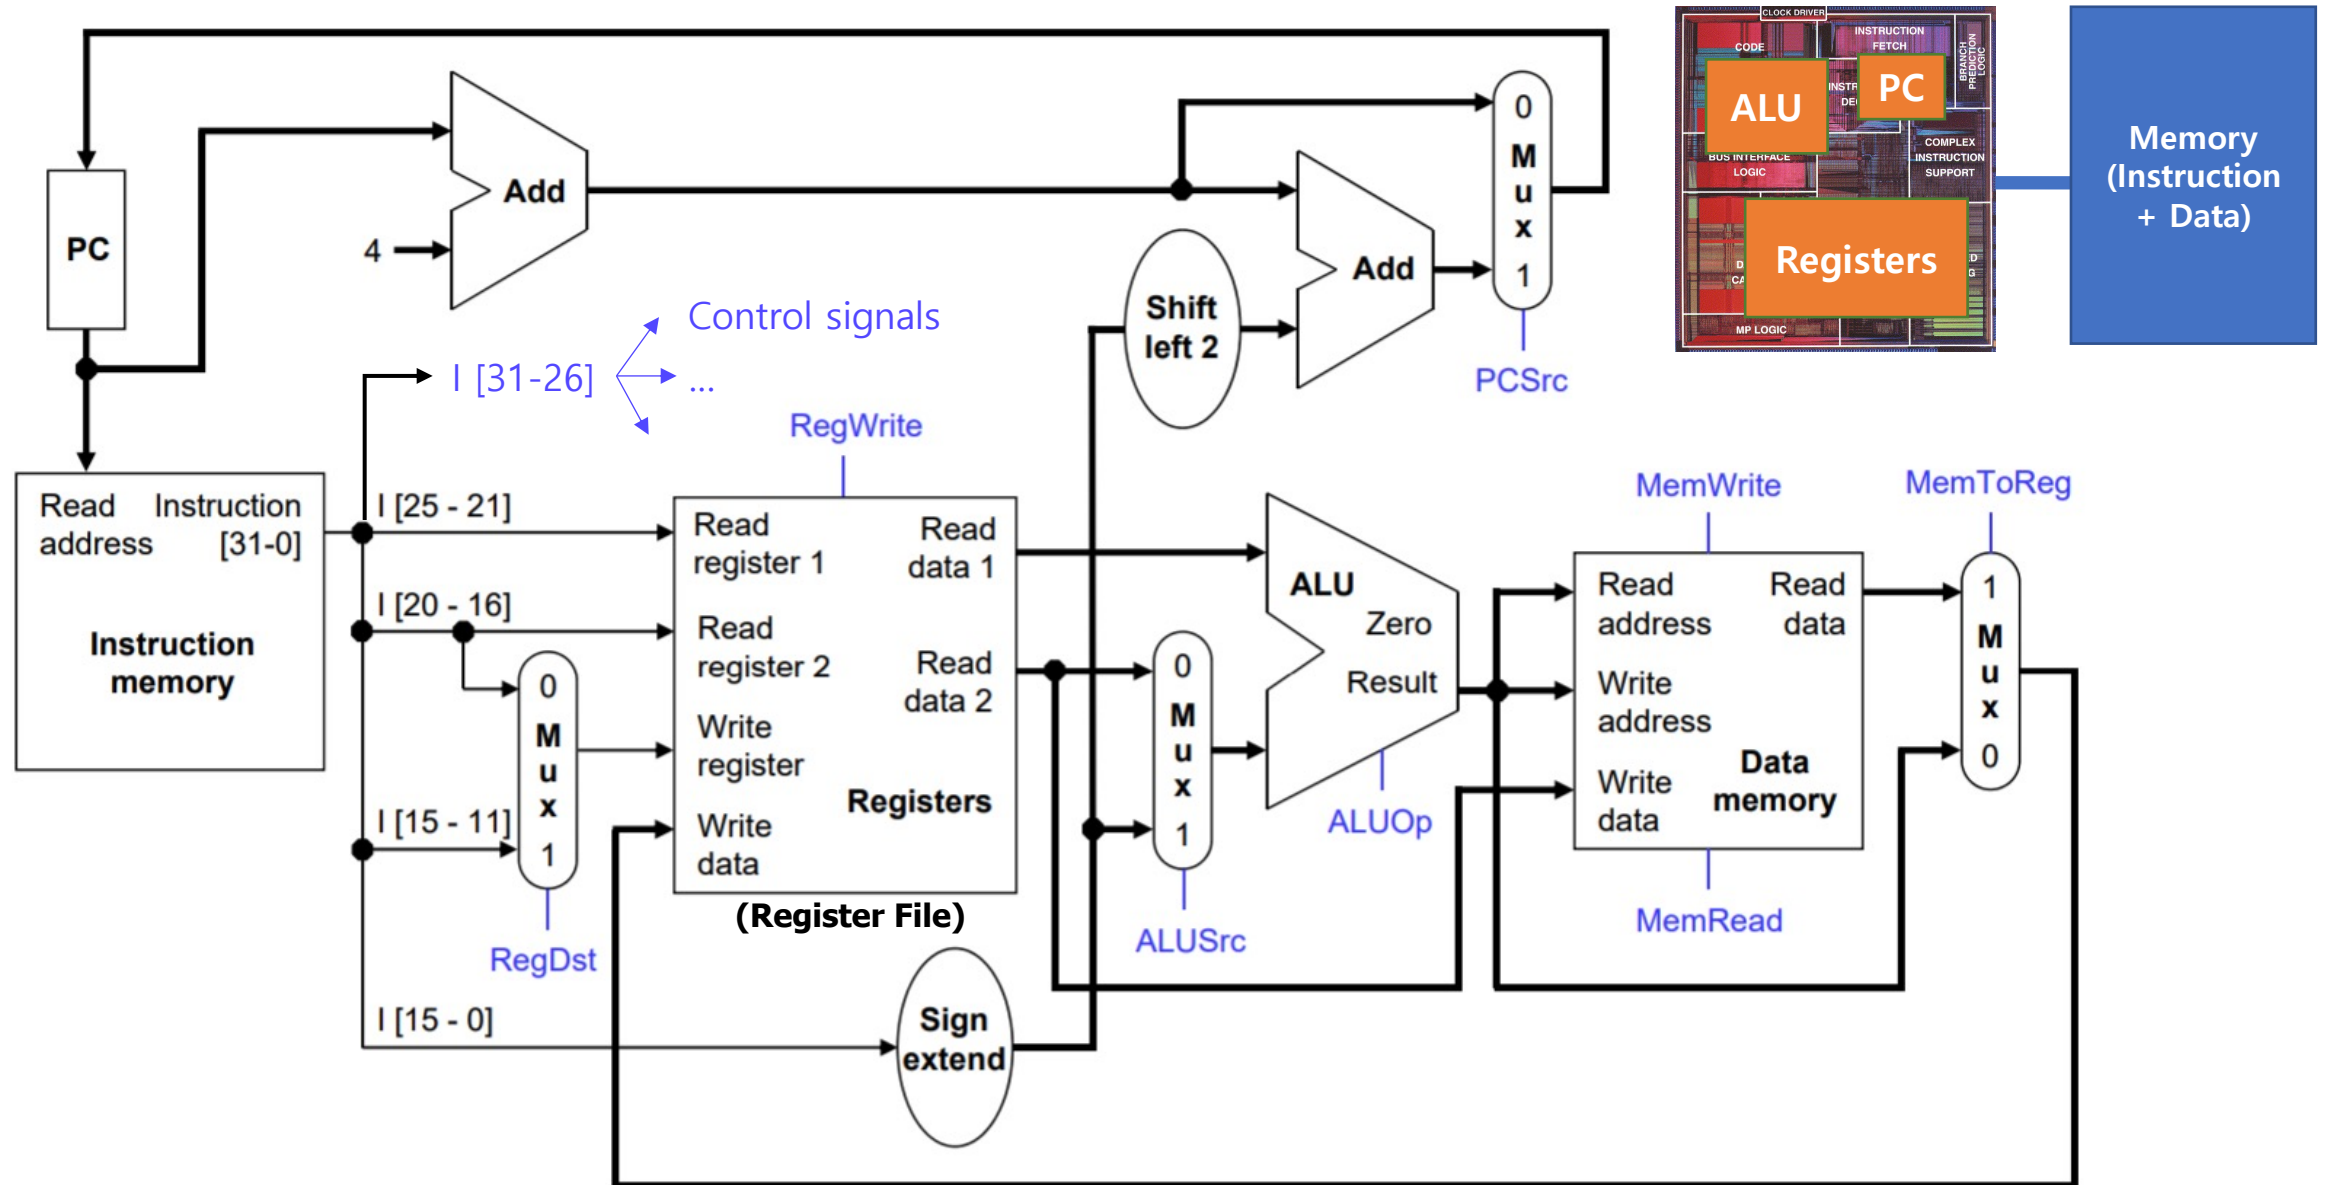


More Complex Module



CPU

Simple CPU with Single Cycle Datapath



Five Stages of Datapath

Instruction
Fetch

IF

Instruction
Decode

ID

Execution

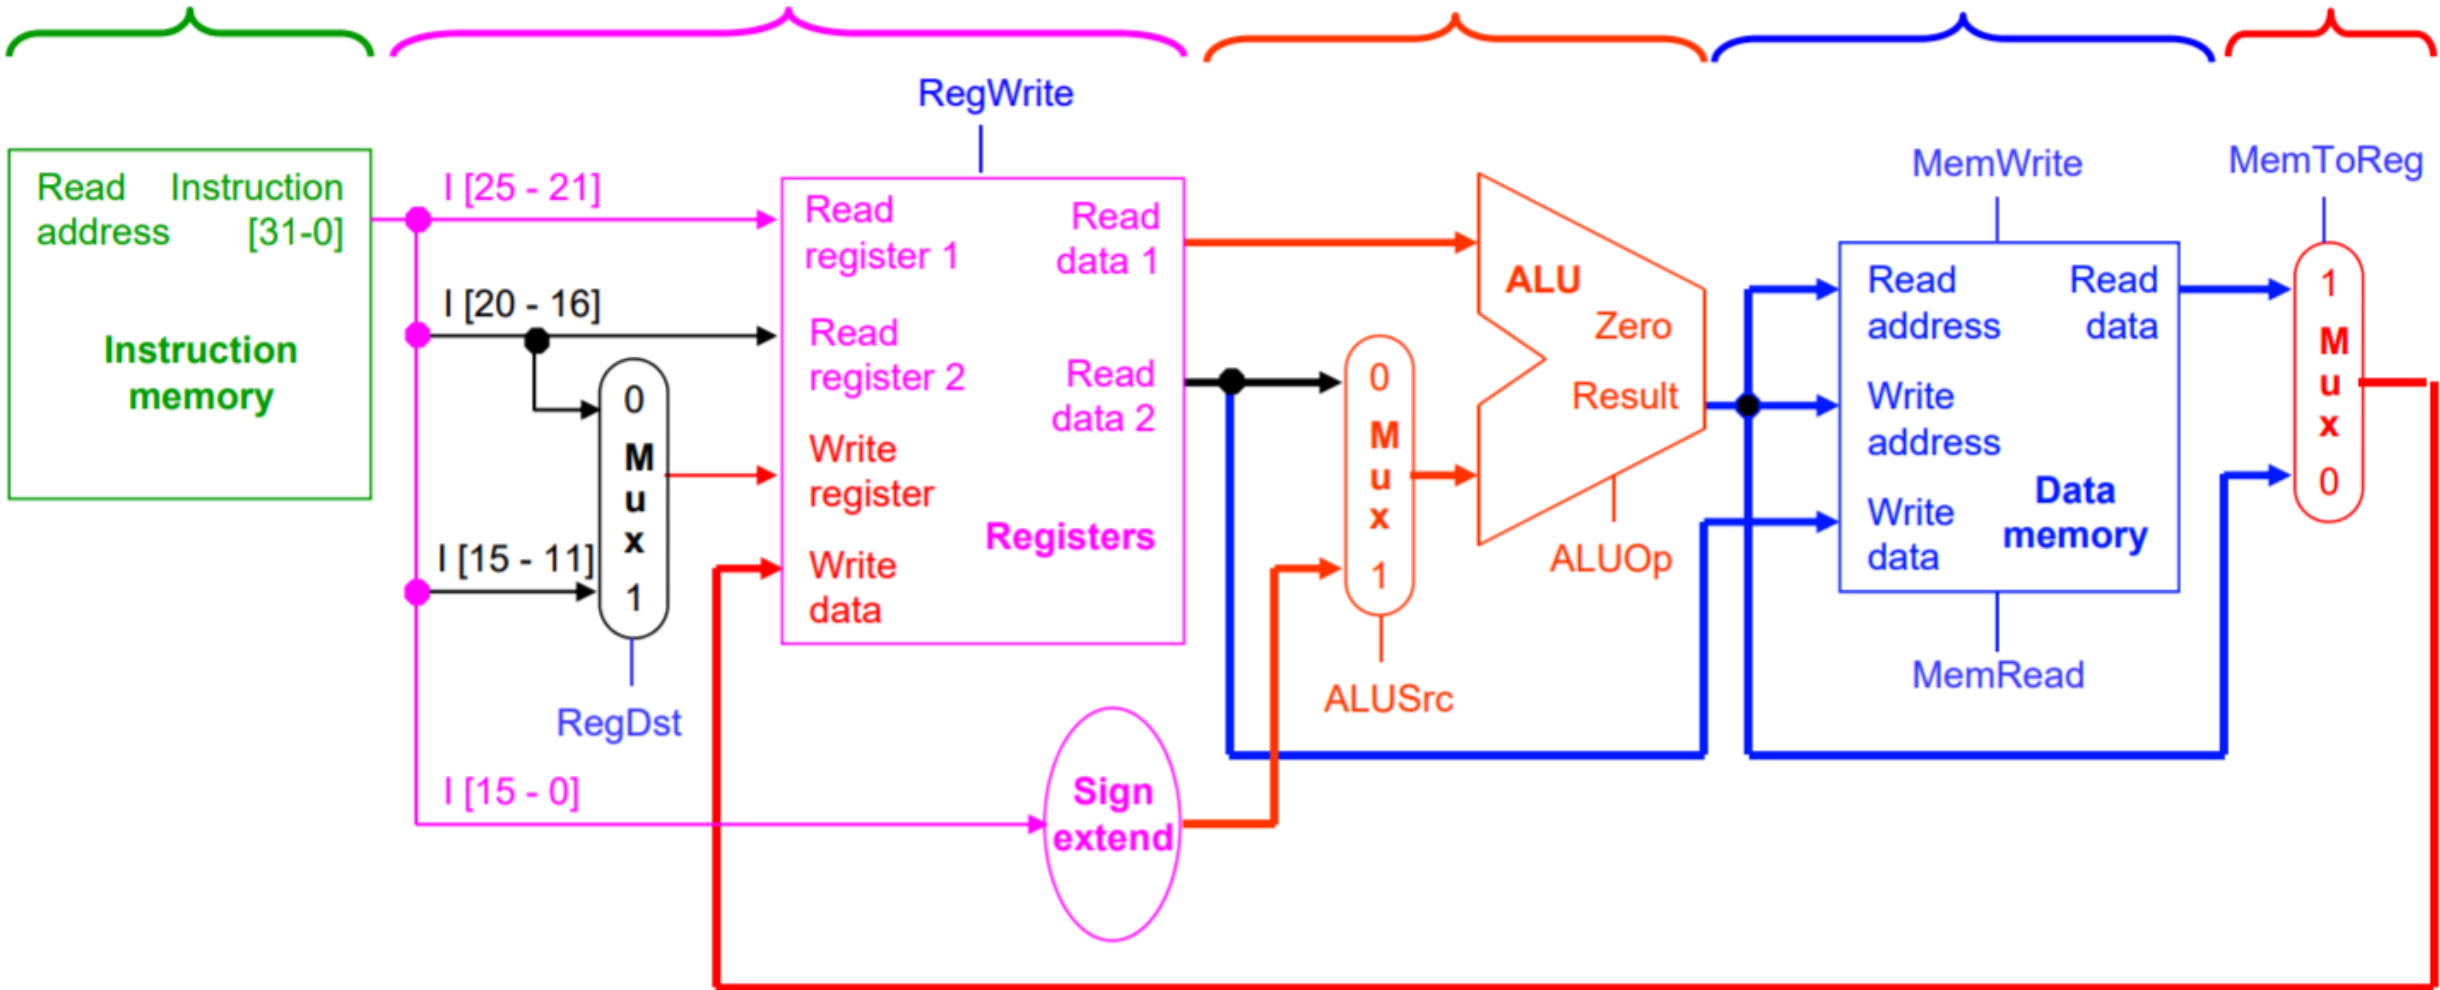
EXE

Memory
Access

MEM

Write
Back

WB




Five Stages of Instruction Execution


- IF (Instruction Fetch)
 - Reads the instruction at PC
 - $PC += 4$ (assuming 4-byte RISC instructions)
- ID (Instruction Decode)
 - Understands the instruction
 - Reads registers
- EX (Execute)
 - Performs the operation
 - Arithmetic/Logical Operations
- MEM (Memory Access)
 - Loads from memory or
 - Stores to memory
- WB (Write Back)
 - Writes the result to appropriate registers

Laundry Pipeline Example

No pipelining: only 1 worker




With pipelining: 4 workers



Pipelining

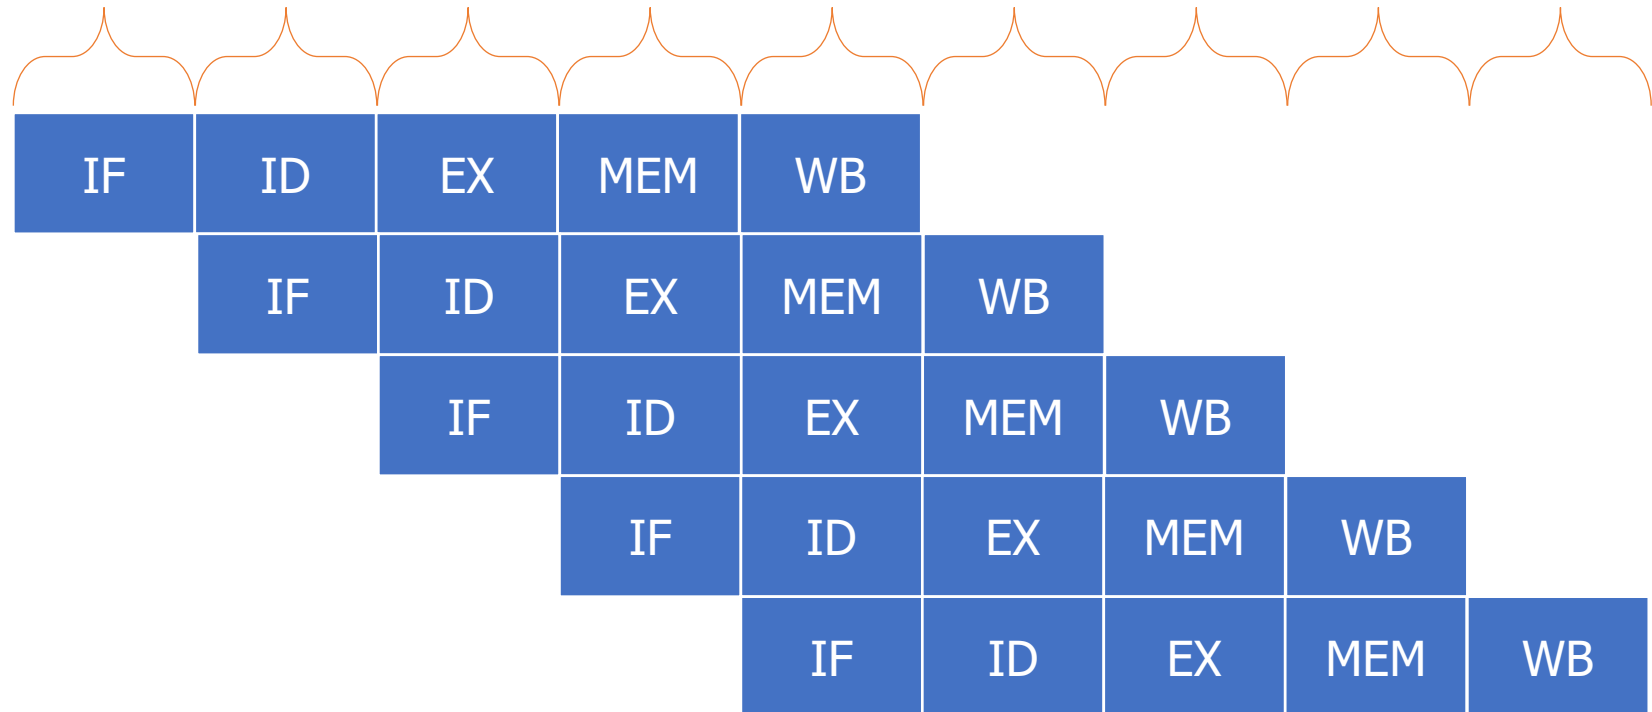
Single Instruction / Single CPU Cycle

**Sequential
Execution**

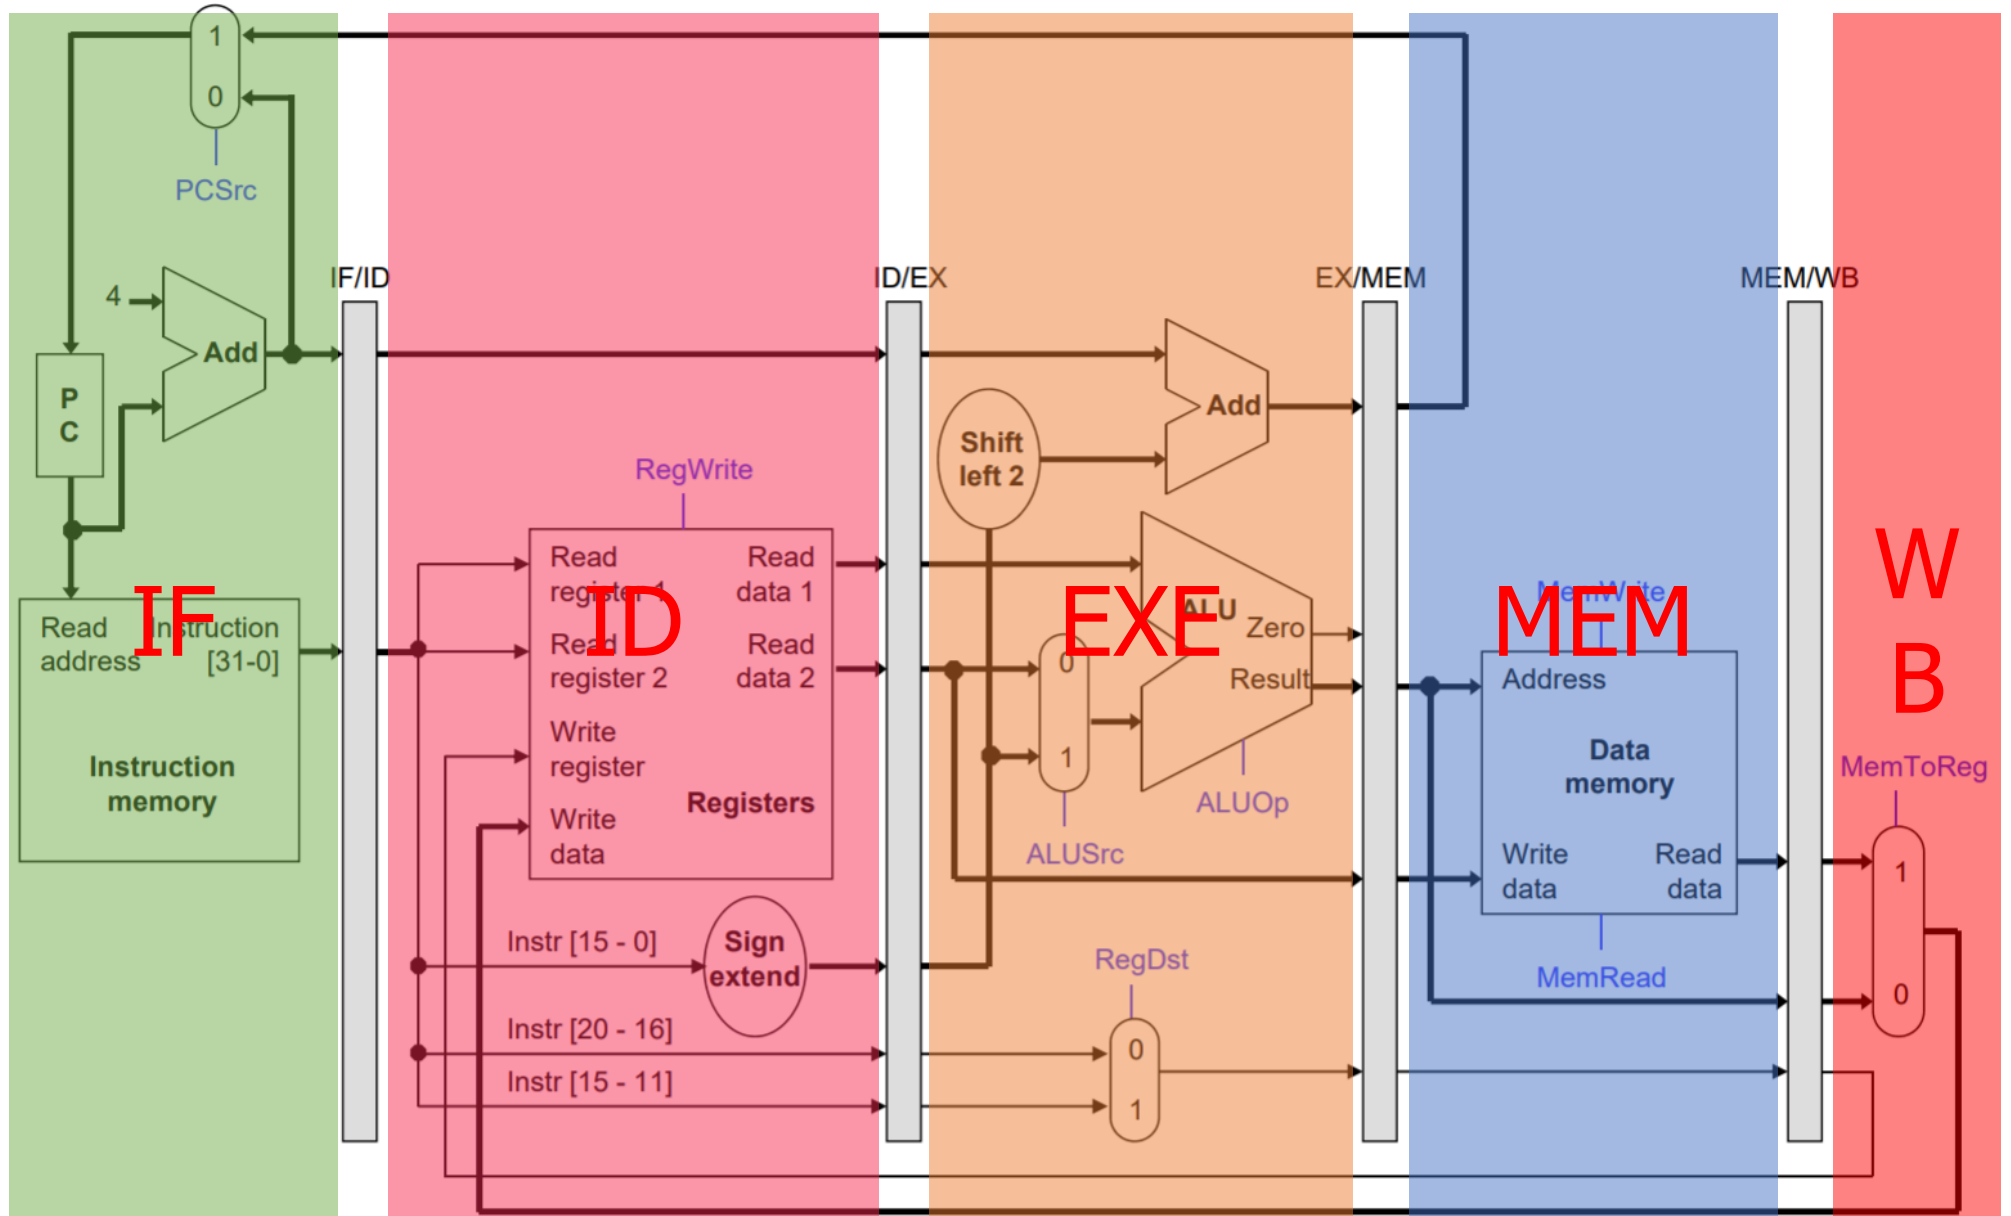


Single Stage / Single CPU Cycle

**(Ideal)
Pipelined
Execution**



Pipelined Datapath

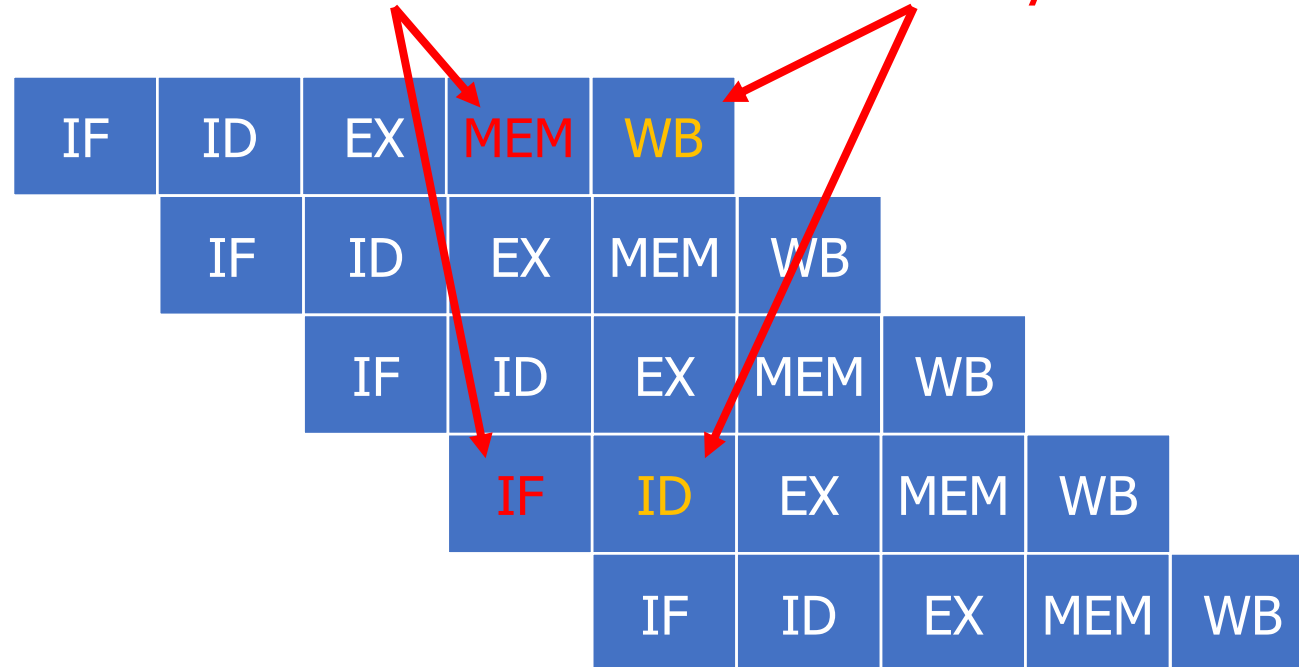


Pipeline Hazards

- Structural Hazards
 - HW resource conflicts
 - Harvard architecture is better in terms of pipelining

Both try to access memory


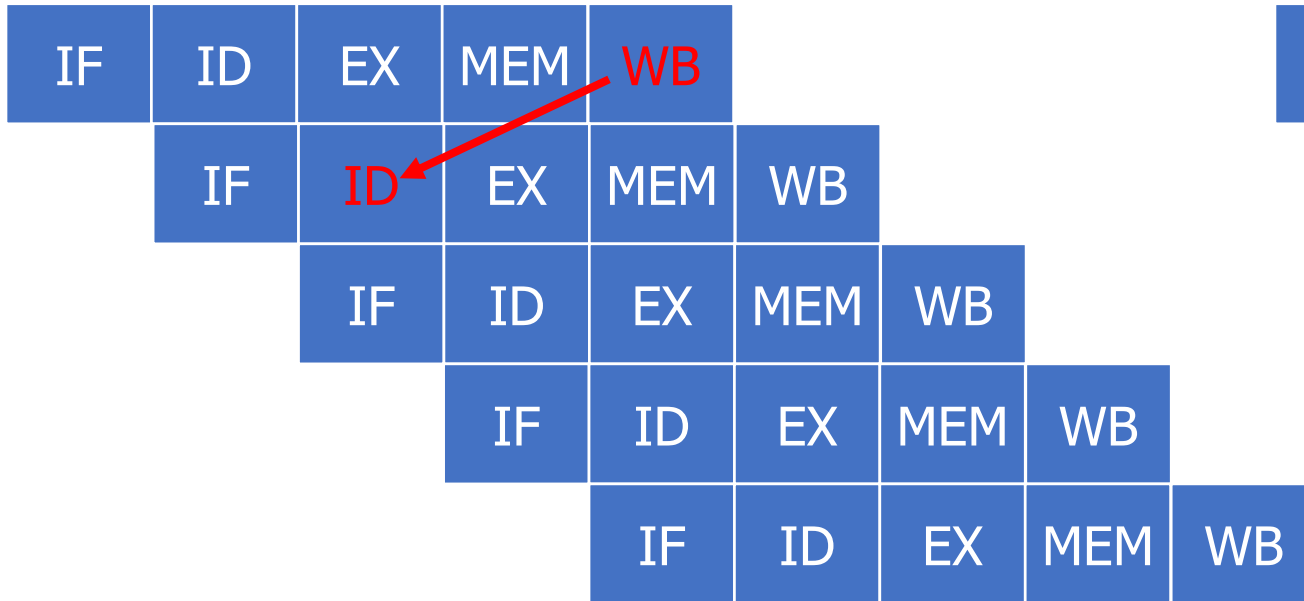
Both try to access register file



Pipeline Hazards

- Data Hazards
 - Data dependencies
 - RAW, WAR, WAW

```
add    r3, r2, r3 # r3 = r2 + r3
mov    r4, r3     # r4 = r3
```



Pipeline Hazards

- Control Hazards
 - Control uncertainty
 - Conditional branch

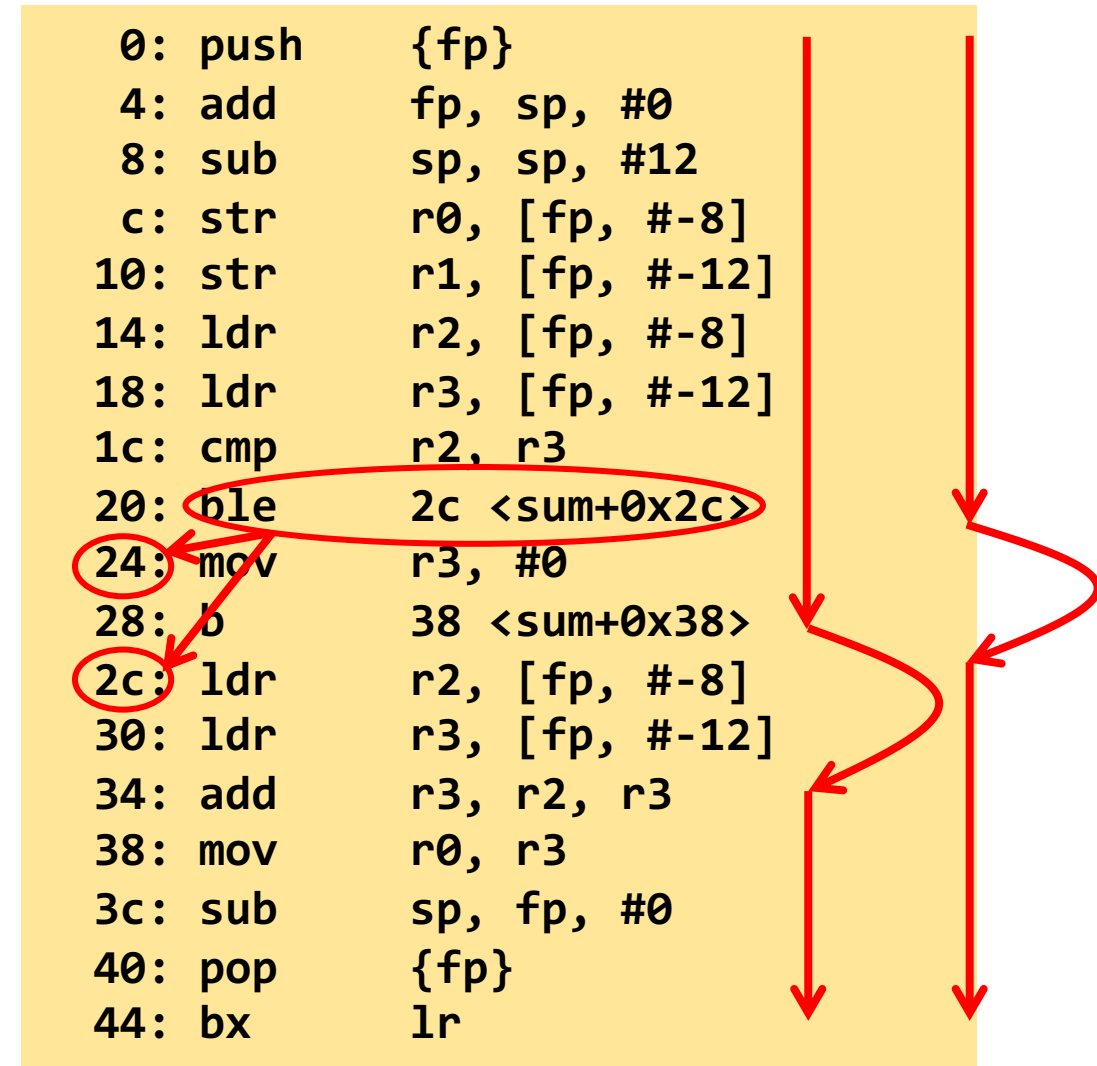
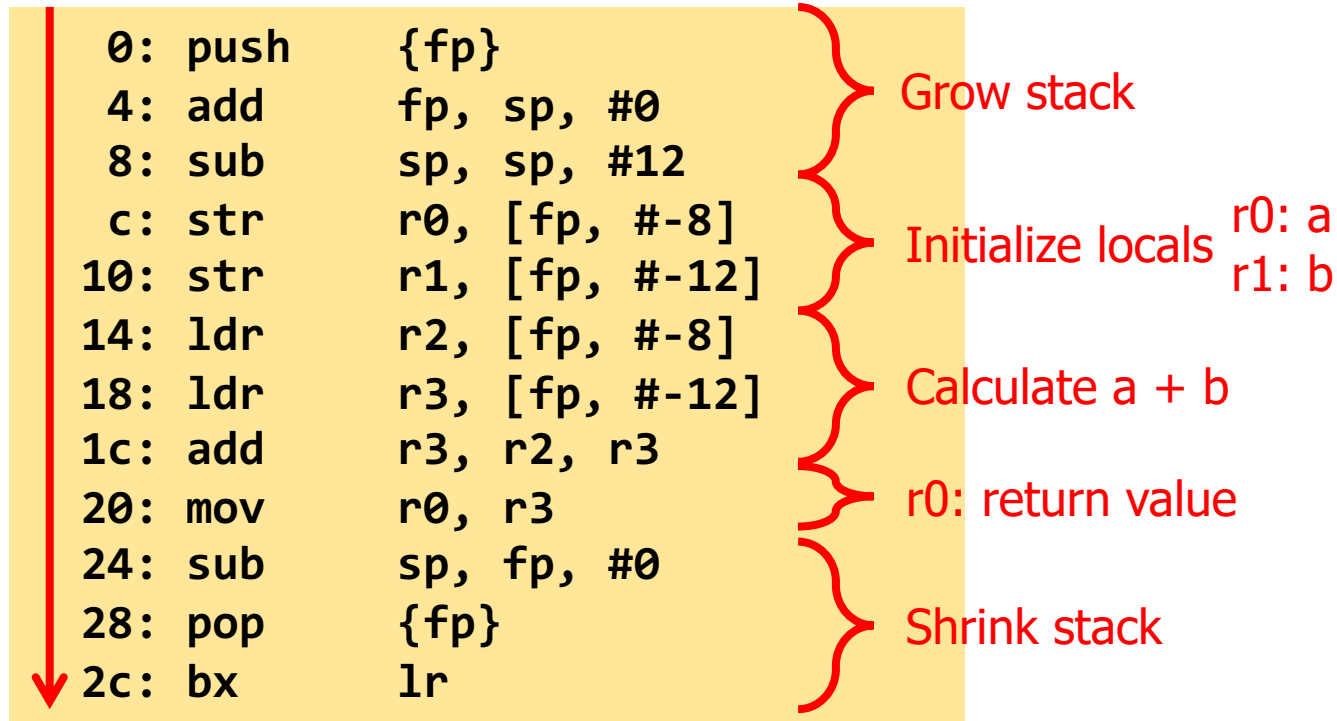
```
int sum(int a, int b)
{
    return a + b;
}
```

- No control uncertainty

```
int sum(int a, int b)
{
    if (a > b) {
        return 0;
    }
    else {
        return a + b;
    }
}
```

- Control uncertainty
 - if or else

Control Uncertainty



*. ble: branch less than or equal


Speculation and Branch Prediction

- Speculative execution
 - Doing (possibly) useless is better than doing nothing
 - In case of incorrect prediction, flush the pipeline and begin from the start
- Branch prediction
 - Improves the probability of taking the right branch
 - Based on the historical data

Out-Of-Order Execution

- Changes the execution order of instructions at runtime

```
1: A = B + C;  
2: E = A + F;  
3: D = B + 1;
```



```
1: A = B + C;  
2: D = B + 1;  
3: E = A + F;
```

- Line 2 depends on Line 1's result
 - High probability of pipeline stalls
- Reordered instructions
 - Less probability of pipeline stalls
 - Produces the same result as the original execution order

Processor Performance Metrics

- Latency (execution time)
 - Time to finish a program
 - A program's perspective
- Throughput (bandwidth)
 - Number of programs processed in a certain time
 - A system's perspective

Taxi is better
for latency

Bus is better
for throughput



Execution Time

- Seconds / cycle
 - Higher CPU clock frequency (no longer possible)
 - Energy consumption and heat problem
 - Multicore CPU can provide more CPU cycles
- Cycles / instruction
 - CISC CPU takes more CPU cycles per instruction (RISC wins)
- Instructions / program
 - RISC compilers produce more instructions for a given program (CISC wins)

$$\frac{\text{seconds}}{\text{program}} = \frac{\text{seconds}}{\text{cycle}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{instructions}}{\text{program}}$$


Throughput

- MIPS (Million Instructions Per Second)
 - How many instructions can a CPU execute in a second
 - Historically important but not that useful these days

$$\text{MIPS} = \frac{\text{instructions}}{\text{cycles}} \times \frac{\text{cycles}}{\text{second}}$$

Will high throughput always lead to low latency?

Same Throughput, Different Latencies



Summary

- Pipeline Architecture
- Pipeline Hazards
 - Structural Hazards
 - Data Hazards
 - Control Hazards
- Performance Metrics
 - Latency
 - Throughput